

Owning your agent systems.

Building digital employees without becoming a software company.

Dan Schnurbusch · Schnurbusch Law

You don't need to own every layer. Understand what you control, what you rent, and who is responsible when something goes wrong.

Buy. Configure. Operate.

Buy a legal-platform agent.

Examples: Filevine LOIS; Clio Work / Vincent. You gain integrated legal workflows and vendor support. You accept product boundaries and platform dependence. A good fit when your recurring job fits the product.

Configure a commercial general-purpose agent.

Examples: Claude / ChatGPT, with capabilities that vary by app, agent surface, and plan. You gain broad capability with less infrastructure to operate. You accept vendor limits and provider dependence.

Operate an open harness.

Examples: Hermes / OpenClaw. You gain code access and model/runtime choice. You accept responsibility for integrations, security, upgrades, and maintenance. A good fit when distinct workflows justify an operating owner.

Conceptual decision map, not a product ranking. Categories overlap. Open source does not automatically mean safer or cheaper. First draft.

The model is only one layer.

Model

Reasons and generates. Can be proprietary or open, local or hosted.

Harness

Manages the work loop, context, tools, and execution.

Tools / MCP

Connects the agent to data and actions. MCP is a connection standard, not a blanket grant of authority.

Skills

Packages procedures, standards, templates, scripts, and references. A skill is a procedure, not permission.

Memory & source records

Memory supports continuity. Canonical records and original documents supply facts.

Permissions & verification

Limits actions and checks results. Enforce access in tools and services, not only in prompts.

Three separate choices.

Software: Which harness manages the job?

Inference: Which model does the thinking, and where?

Hosting: Where does the harness run?

Open harness ≠ open model.

Self-hosted harness ≠ all data stays local.

A chat connection or scheduled trigger starts the job; it does not expand its authority. Map data destinations, provider terms, retention, and access before connecting confidential information.

Digital employee is an operating metaphor. Human professional judgment and accountability remain necessary.

Start with the assignment.

Review matters waiting on clients. Check whether the missing information already arrived. Prepare the next action for my review.

Trigger → retrieve → check → apply procedure
→ draft → human approval → verify

Illustrative example

The current record asks for a signed form and a supporting statement. A recent email includes the signed form. The agent reads the attachment and confirms it is the right form. It prepares a short, unsent request for only the statement, with source references and a review link.

It must not: send the message, treat an attachment filename as reviewed contents, or mark the entire request complete.

Design for the failure

- **Stale evidence:** compare current records, message dates, and actual attachments.
- **False completion:** read back the exact target. An API success is not proof that the intended draft or update exists.
- **Missing information:** preserve unknowns and stop when identity or authority is unclear.
- **Uncertain write:** investigate the destination before retrying. Repeating an action can create duplicates.

A skill is a short procedure.

1. Read the current record.
2. Check the latest correspondence and attachments.
3. Identify only what remains missing.
4. Prepare an unsent draft.
5. Return the source references and review link.

All examples are fictional. Nothing here is a live client record or a measured performance result.

Share procedures. Separate identities.

- Use named users and least-privilege access.
- Limit access by role and matter.
- Separate read, draft, and send permissions.
- Require approval for the exact consequential action.
- Name an owner for access, updates, failures, and costs.
- Keep useful audit records without needlessly copying client content.

Cost follows the work.

Retrieve relevant fields, not every record. Load the needed procedure, not the entire library. Use ordinary code for predictable filtering, extraction, and arithmetic. Reserve strong model reasoning for judgment that warrants it.

Cost per accepted result

(model + infrastructure + review + maintenance)

÷ accepted results

Measure reviewer time and quality together. A cheap draft that takes longer to fix is not a cheap result. Local inference can exchange provider charges for hardware, electricity, maintenance, and different capabilities.

Start with a controlled pilot.

Use nonconfidential or approved sample data. Include an old message, a missing document, an ambiguous identity, and a failed save. Measure accepted results, review time, errors, and operating cost before granting more authority.

No savings figures or benchmark scores are claimed. Product features and terms change; verify the exact product and plan.

Write the job description.

Job & deliverable

What useful artifact should exist? Who or what starts the job?

Sources

Which canonical records and documents supply the facts?

Procedure & stop conditions

What steps apply? What uncertainty makes the agent stop?

Allowed actions & approval

What can it read, draft, or change? Who approves each consequential action?

Proof of completion

Which exact target must be read back or inspected?

Data routing & operating owner

Where does information go? Who maintains the system?

Pilot & success measure

Which normal and failure cases will you test? How will you measure cost and quality?

Own the procedure. Choose what else to own.

Then select the system that can support the job.

Go deeper.

Filevine LOIS Console

filevine.com/news/filevine-launches-lois-console-ai-that-runs-the-firm

Clio Work / Vincent

help.clio.com/hc/en-us/articles/48791655710235-Understand-Clio-Work

Claude Enterprise

claude.com/solutions/enterprise

Claude custom MCP connectors

support.claude.com/en/articles/11176164-use-connectors-to-extend-claude-s-capabilities

Hermes Agent documentation

hermes-agent.nousresearch.com/docs/

ChatGPT connected apps

help.openai.com/en/articles/11487775-connected-apps-in-chatgpt

OpenClaw tools & skills

docs.openclaw.ai/tools

Model Context Protocol

modelcontextprotocol.io

Agent Skills

agentskills.io

Vendor sources establish product descriptions, not independently verified performance. The continuum describes architectural tradeoffs. It is not a security certification, procurement recommendation, or benchmark.

Keep this PDF if the hosted site is later removed. The HTML presentation also works offline when saved.